

## Simple Linear Regression Assignment

Data Set : Salary\_Data Building a Prediction Model for Salary Hike

### 1. Import Necessary libraries

```
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns

import statsmodels.formula.api as smf

from sklearn.metrics import mean_squared_error
from math import sqrt

import warnings
warnings.filterwarnings('ignore')
```

#### 1. Import Data

```
salary_details = pd.read_csv('/content/Salary_Data (2).csv')
salary_details
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0

```

25          9.0  105582.0
26          9.5  116969.0
27          9.6  112635.0
28         10.3  122391.0
29         10.5  121872.0

```

## 1. Data Understanding

a) Initial Analysis:

```
salary_details.head()
```

```

   YearsExperience  Salary
0             1.1  39343.0
1             1.3  46205.0
2             1.5  37731.0
3             2.0  43525.0
4             2.2  39891.0

```

```
salary_details.shape
```

```
(30, 2)
```

```
salary_details.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes

```

```
salary_details.isna().sum()
```

```

YearsExperience    0
Salary             0
dtype: int64

```

b) Correlation Matrix :

```

corr_matrix = salary_details.corr()
corr_matrix

```

```

           YearsExperience  Salary
YearsExperience    1.000000  0.978242
Salary             0.978242  1.000000

```

```

sns.heatmap(data = corr_matrix,annot = True)
plt.show()

```




## 1. Perform Assumption Check

### a) Outlier Test Using Box Plot :

```
plt.figure(figsize = (12,5))
plt.subplot(1,2,1)
salary_details['YearsExperience'].hist()
plt.subplot(1,2,2)
salary_details.boxplot(column = ['YearsExperience'])
```


```
plt.show()
```



```
plt.figure(figsize = (12,5))
plt.subplot(1,2,1)
salary_details['Salary'].hist()
```

```
plt.subplot(1,2,2)
salary_details.boxplot(column = ['Salary'])
```


```
plt.show()
```




From the above histograms and boxplots, we found that there is no outliers present inside the YearsExperience and Salary data.

#### **b) Normality / Distribution Test Using Distplot :**

```
sns.distplot(salary_details['YearsExperience'])
plt.show()
```



```
sns.distplot(salary_details['Salary'])  
plt.show()
```



Normality Test Failed

## 7. Model Building || 8. Model Training

Now Try to fit Model for Salary Hike

### Model 1: Without Applying any Transformation

Using Statsmodel

```
model_1 = smf.ols(formula = 'YearsExperience~Salary', data =  
salary_details).fit()  
model_1
```

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at  
0x7fdcb5b4bdf0>
```

```
#coefficient  
model_1.params
```

```
Intercept    -2.383161  
Salary         0.000101  
dtype: float64
```

```
model_1.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

### OLS Regression Results

```
=====
Dep. Variable:          YearsExperience    R-squared:
0.957
Model:                  OLS              Adj. R-squared:
0.955
Method:                Least Squares     F-statistic:
622.5
Date:                  Fri, 20 Jan 2023   Prob (F-statistic):
1.14e-20
Time:                  17:28:01          Log-Likelihood:
-26.168
No. Observations:      30               AIC:
56.34
Df Residuals:          28               BIC:
59.14
Df Model:              1

Covariance Type:      nonrobust
```

```
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept    -2.3832      0.327     -7.281      0.000     -3.054
-1.713
Salary        0.0001    4.06e-06     24.950      0.000      9.3e-05
0.000
=====
```

```
=====
Omnibus:      3.544    Durbin-Watson:
1.587
Prob(Omnibus): 0.170    Jarque-Bera (JB):
2.094
Skew:         -0.412    Prob(JB):
0.351
Kurtosis:     2.003    Cond. No.
2.41e+05
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[2] The condition number is large, 2.41e+05. This might indicate that
there are
strong multicollinearity or other numerical problems.
"""
```

From the Above OLS Regression Result the R-Squared value is  $0.957 > 0.75$  and we can say that this Model is good to Predict Salary\_hike and p-value  $< 0.05$  and it is significant model

## Model 2: Apply Log Transformation of Y

```
model_2 = smf.ols(formula = 'Salary~np.log(YearsExperience)',data =
salary_details).fit()
model_2
```

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at
0x7fdcb5b72250>
```

```
model_2.params
```

```
Intercept          14927.97177
np.log(YearsExperience)  40581.98796
dtype: float64
```

```
model_2.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          Salary    R-squared:
0.854
Model:                  OLS      Adj. R-squared:
0.849
Method:                 Least Squares    F-statistic:
163.6
Date:                   Fri, 20 Jan 2023    Prob (F-statistic):
3.25e-13
Time:                   17:29:24    Log-Likelihood:
-319.77
No. Observations:      30    AIC:
643.5
Df Residuals:          28    BIC:
646.3
Df Model:              1

Covariance Type:      nonrobust

=====
=====
```

	coef	std err	t	P> t
[0.025      0.975]				
-----				
Intercept	1.493e+04	5156.226	2.895	0.007
4365.921      2.55e+04				
np.log(YearsExperience)	4.058e+04	3172.453	12.792	0.000
3.41e+04      4.71e+04				
=====				
Omnibus:	1.094	Durbin-Watson:		
0.512				
Prob(Omnibus):	0.579	Jarque-Bera (JB):		
0.908				
Skew:	0.156	Prob(JB):		
0.635				
Kurtosis:	2.207	Cond. No.		
5.76				
=====				
=====				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
 """

### Model 3: Apply Log Transformation of X

```
model_3 = smf.ols(formula = 'np.log(Salary)~YearsExperience',data = salary_details).fit()
model_3
```

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at 0x7fdccef66940>
```

```
model_3.params
```

```
Intercept            10.507402
YearsExperience       0.125453
dtype: float64
```

```
model_3.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:            np.log(Salary)    R-squared:
0.932
```



```

Model:                                OLS    Adj. R-squared:
0.930
Method:                            Least Squares    F-statistic:
383.6
Date:                            Fri, 20 Jan 2023    Prob (F-statistic):
7.03e-18
Time:                            17:30:10    Log-Likelihood:
28.183
No. Observations:                    30    AIC:
-52.37
Df Residuals:                        28    BIC:
-49.56
Df Model:                            1

```

Covariance Type: nonrobust

```

=====
=====
                                coef    std err          t      P>|t|
[0.025    0.975]
-----
-----
Intercept                10.5074      0.038    273.327      0.000
10.429    10.586
YearsExperience           0.1255      0.006    19.585      0.000
0.112    0.139
=====
=====
Omnibus:                  0.826    Durbin-Watson:
1.438
Prob(Omnibus):            0.661    Jarque-Bera (JB):
0.812
Skew:                    0.187    Prob(JB):
0.666
Kurtosis:                2.286    Cond. No.
13.2
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

#### Model 4: Apply Log Transformation of X and Y

```

model_4 = smf.ols(formula =
'np.log(Salary)~np.log(YearsExperience)',data = salary_details).fit()
model_4

```

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at 0x7fdcb5b6b9a0>
```

```
model_4.params
```

```
Intercept          10.328043
np.log(YearsExperience)  0.562089
dtype: float64
```

```
model_4.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

### OLS Regression Results

```
=====
```

```
=====
```

```
Dep. Variable:          np.log(Salary)    R-squared:
0.905
Model:                  OLS              Adj. R-squared:
0.902
Method:                 Least Squares     F-statistic:
267.4
Date:                  Fri, 20 Jan 2023   Prob (F-statistic):
7.40e-16
Time:                  17:30:56          Log-Likelihood:
23.209
No. Observations:      30               AIC:
-42.42
Df Residuals:          28               BIC:
-39.61
Df Model:              1
```

```
Covariance Type:      nonrobust
```

```
=====
```

```
=====
```

		coef	std err	t	P> t
--	--	------	---------	---	------

[0.025	0.975]				
--------	--------	--	--	--	--

```
-----
```

```
-----
```

Intercept		10.3280	0.056	184.868	0.000
10.214	10.442				
np.log(YearsExperience)		0.5621	0.034	16.353	0.000
0.492	0.632				

```
=====
```

```
=====
```

```
Omnibus:              0.102    Durbin-Watson:
0.988
Prob(Omnibus):         0.950    Jarque-Bera (JB):
0.297
```

```
Skew:                0.093    Prob(JB):
0.862
Kurtosis:            2.549    Cond. No.
5.76
```

```
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

### Model 5: Apply Square Root Transformation

```
model_5 = smf.ols(formula = 'Salary~np.sqrt(YearsExperience)',data =
salary_details).fit()
model_5
```

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at
0x7fdcb5aef4f0>
```

```
model_5.params
```

```
Intercept                -16055.769117
np.sqrt(YearsExperience)  41500.680583
dtype: float64
```

```
model_5.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

#### OLS Regression Results

```
=====
=====
```

```
Dep. Variable:            Salary    R-squared:
0.931
Model:                    OLS      Adj. R-squared:
0.929
Method:                   Least Squares    F-statistic:
377.8
Date:                     Fri, 20 Jan 2023    Prob (F-statistic):
8.57e-18
Time:                     17:33:21    Log-Likelihood:
-308.52
No. Observations:         30    AIC:
621.0
Df Residuals:             28    BIC:
623.8
Df Model:                  1
```

Covariance Type: nonrobust

```
=====
=====
              coef      std err          t      P>|t|
[0.025      0.975]
-----
Intercept      -1.606e+04    4921.599     -3.262     0.003
-2.61e+04    -5974.331
np.sqrt(YearsExperience)  4.15e+04    2135.122     19.437     0.000
3.71e+04     4.59e+04
=====
=====
Omnibus:                0.588    Durbin-Watson:
1.031
Prob(Omnibus):          0.745    Jarque-Bera (JB):
0.638
Skew:                   0.011    Prob(JB):
0.727
Kurtosis:               2.286    Cond. No.
9.97
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
"""

**CONCLUSION** = Comparing between all Models we got to know that without applying any transformation for the Model\_1 we got the Higher R-squared Value i.e. 0.957 as compared to all Model.

Hence the Model\_1 is better model to predict Salary\_hike